
MLOps Group Assignment

AG News Text Classification Pipeline

Fine-tuning DistilBERT with Docker, CI/CD, and W&B Tracking

Course: MLOps — PGD AI Program
Institute: IIT Jodhpur
Group Number: Group 13
Submission Date: June 2026

Name	Roll Number	Contribution
Ritesh Maury	G25AIT2087	GitHub Setup, Project Lead, Merges
Kosuru Yuvaraj	G25AIT2054	Training, Docker, CI/CD, HF Hub, Report
Laishram Khumanleima Chanu	G25AIT2056	Data Preparation, Inference Script
Guntakal Sandeep Kumar	G25AIT2036	W&B Tracking, Evaluation, Testing

All Submission Links

GitHub Repository: [riteshmaury-iitj/group13-assignment-mlops](https://github.com/riteshmaury-iitj/group13-assignment-mlops)
Kaggle Notebook V1: [kyuvarajg25ait2054/group13-ag-news-k](https://kaggle.com/kyuvarajg25ait2054/group13-ag-news-k)
Kaggle Notebook V2: [kyuvarajg25ait2054/group13-ag-news-k](https://kaggle.com/kyuvarajg25ait2054/group13-ag-news-k)
HuggingFace Model: [YuvarajK-g25ait2054/ag-news-distilbert](https://huggingface.co/YuvarajK-g25ait2054/ag-news-distilbert)
Docker Image: [yuvarajkg25ait2054/ag-news-classifier](https://hub.docker.com/r/yuvarajkg25ait2054/ag-news-classifier)
W&B Dashboard: [g25ait2054-iit-jodhpur/mlops-assignment3](https://wandb.ai/g25ait2054-iit-jodhpur/mlops-assignment3)

Member Profiles & Project Links

All links are public and live at submission time.

Member Profiles

Member	GitHub	Kaggle	HuggingFace	W&B
Ritesh Maury G25AIT2087	riteshmaury-iitj	ritesh329	Recurrent	models-iit-jodhpur
Kosuru Yuvaraj G25AIT2054	yuvarajkosuru	kyuvarajg25ait2054	YuvarajK-g25ait2054	g25ait2054
Laishram K. Chanu G25AIT2056	group13-mlops	khumanleima	LKHUMANLEIMA	g25ait2056
Guntakal Sandeep G25AIT2036	Sandeep_dev	guntakalsandeep	guntaksa	g25ait2036

Individual Member Links

Roll No	Resource	Link
G25AIT2087	GitHub	https://github.com/riteshmaury-iitj
	Prepared Data (HF)	https://huggingface.co/datasets/Recurrent/prepared_data_mlops2
	Data Preparation	https://github.com/riteshmaury-iitj/group13-assignment-mlops/blob/main/notebooks/prepare_data.ipynb
	Model Training (Kaggle)	https://www.kaggle.com/code/ritesh329/train-model
	Model Inference (Kaggle)	https://www.kaggle.com/code/ritesh329/model-inferenence
	HuggingFace Model	https://huggingface.co/Recurrent/ag-news-distilbert
	W&B	https://wandb.ai/models-iit-jodhpur/mlops-assignment-group13
G25AIT2054	GitHub	https://github.com/yuvarajkosuru
	Kaggle Notebook	https://www.kaggle.com/code/kyuvarajg25ait2054/group13-ag-news-k
	HuggingFace Model	https://huggingface.co/YuvarajK-g25ait2054/ag-news-distilbert
	W&B	https://wandb.ai/g25ait2054-iit-jodhpur/mlops-assignment3
G25AIT2056	GitHub	https://github.com/riteshmaury-iitj/group13-assignment-mlops
	Kaggle Notebook	https://www.kaggle.com/code/khumanleima/notebook35f6cea6a6
	HuggingFace Model	https://huggingface.co/LKHUMANLEIMA/ag-news-distilbert
	W&B	https://wandb.ai/g25ait2056-indian-institute-of-technology/mlops-assignment3
G25AIT2036	GitHub (branch)	https://github.com/riteshmaury-iitj/group13-assignment-mlops/tree/Sandeep_dev
	Kaggle Notebook	https://www.kaggle.com/code/guntakalsandeepkumar/g25ait2036
	HuggingFace Model	https://huggingface.co/guntaksa/AG-News-distilbert1
	W&B	https://wandb.ai/g25ait2036-indian-institutue-of-technology-mlops-assignment3

Project Links (all public)

Component	Link
GitHub Repository (public)	github.com/riteshmaury-iitj/group13-assignment-mlops
Kaggle Notebook — Version 1 (public)	kaggle.com/code/kyuvarajg25ait2054/group13-ag-news-k
Kaggle Notebook — Version 2 (public)	kaggle.com/code/kyuvarajg25ait2054/group13-ag-news-k
HuggingFace Model (public)	huggingface.co/YuvarajK-g25ait2054/ag-news-distilbert
Docker Image (publicly pullable)	hub.docker.com/r/yuvarajkg25ait2054/ag-news-classifier
W&B Project Dashboard (public)	wandb.ai/g25ait2054-iit-jodhpur/mlops-assignment3

1. GitHub Repository, Data Preparation and Model Selection (Tasks 1–3)

1.1 Task 1 — GitHub Repository Setup [10 marks]

Repository [riteshmaurya-iitj/group13-assignment-mlops](https://github.com/riteshmaurya-iitj/group13-assignment-mlops) is public with three branches: **main** (protected, PRs only), **develop** (integration), and **feature/*** (short-lived). Branch protection on **main** requires at least one approval before merging. All work was merged via Pull Request from **develop**.




Figure 1: GitHub repository — public visibility and branch list




Figure 2: GitHub Settings — Branch protection rule and repo details




Figure 3: GitHub Settings — Collaborators: g25ait2036, g25ait2056-beep, YuvarajK-g25ait2054

1.2 Task 2 — Data Preparation & Normalisation [15 marks]

Dataset: `ag_news` from HuggingFace (120,000 train / 7,600 test, 4 balanced classes). Preprocessing applied in `prepare_data.ipynb`:

- Lowercase, remove URLs, strip HTML tags, remove punctuation, normalize whitespace.
- Label mapping saved to `id2label.json`: {0: World, 1: Sports, 2: Business, 3: Sci/Tech}.
- Tokenized with `max_length=128`.

screenshots/label_dist.png

Figure 4: Test set label distribution — 4 balanced classes, 1,900 samples each (7,600 total)




Figure 5: Test set text length distribution — word count per class (mean: 64 words, max: 200)

1.3 Task 3 — Model Selection [10 marks]

Selected **distilbert-base-uncased**: 66M parameters, retains ~97% of BERT performance, and trains on AG News in ~40 min on a T4 GPU. Loaded via `AutoModelForSequenceClassification` with `num_labels=4` and `id2label/label2id` mappings. Tokenizer uses `truncation=True`, `padding=True`, `max_length=128`.

2. Training, W&B Tracking and HuggingFace Deployment (Tasks 4–5)

2.1 Task 4 — Training on Kaggle with W&B [25 marks]

Training on Kaggle (GPU T4 ×2), tracked via W&B project `mlops-assignment3`. Secrets (W&B, HuggingFace, GitHub) stored as Kaggle Secrets, loaded via `UserSecretsClient`.

Hyperparameter Comparison:

Hyperparameter	Version 1 (run-v1)	Version 2 (run-v2)
Base Model	distilbert-base-uncased	distilbert-base-uncased
Learning Rate	2e-5	5e-5
Batch Size (train)	16	64
Epochs	3	5
Weight Decay	0.01	0.01
Max Sequence Length	128	128
Evaluation Strategy	Per epoch	Per epoch
Best Model Metric	F1 (weighted)	F1 (weighted)

Results:

Version	Accuracy	F1 (weighted)	Eval Loss
Version 1 (run-v1)	94.34%	0.9434	0.3679
Version 2 (run-v2)	94.12%	0.9412	0.3575

Version 1 achieved higher accuracy and F1; selected as the final model. Both versions exceeded 94% accuracy.

← Viewing Version 5: 📌 Save & Run All • June 8, 2026 at 2:19 PM

```
In [10]: # Evaluate v1
results_v1 = trainer_v1.evaluate()
print(f"\n=== Version 1 Results ===")
print(f" Accuracy: {results_v1['eval_accuracy']:.4f}")
print(f" F1 Score: {results_v1['eval_f1']:.4f}")
print(f" Eval Loss: {results_v1['eval_loss']:.4f}")

# Finish W&B run
wandb.finish()
print("\nW&B run-ver1 finished.")


/usr/local/lib/python3.12/dist-packages/torch/autograd/function.py:583: UserWarning: Was a
sked to gather along dimension 0, but all input tensors were scalars; will instead unsquee
ze and return a vector.
  return super().apply(*args, **kwargs) # type: ignore[misc]

[119/119 00:14]

wandb: updating run metadata

=== Version 1 Results ===
Accuracy: 0.9434
F1 Score: 0.9434
Eval Loss: 0.3679
```

Figure 6: Kaggle notebook — Version 1 training output (eval accuracy 94.34%, F1 0.9434)



```

In [13]: # Evaluate v2
results_v2 = trainer_v2.evaluate()
print(f"==== Version 2 Results ===")
print(f" Accuracy: {results_v2['eval_accuracy']:.4f}")
print(f" F1 Score: {results_v2['eval_f1']:.4f}")
print(f" Eval Loss: {results_v2['eval_loss']:.4f}")

# Finish W&B run
wandb.finish()
print("\nW&B run-ver2 finished.")

/usr/local/lib/python3.12/dist-packages/torch/autograd/function.py:583: UserWarning: Was a
sked to gather along dimension 0, but all input tensors were scalars; will instead unsquee
ze and return a vector.
  return super().apply(*args, **kwargs) # type: ignore[misc]

[60/60 00:14]

wandb: updating run metadata

==== Version 2 Results ====
Accuracy: 0.9412
F1 Score: 0.9412
Eval Loss: 0.3575

```

Figure 7: Kaggle notebook — Version 2 training output (eval accuracy 94.12%, F1 0.9412)





Figure 8: W&B dashboard — mlops-assignment3 showing run-v1 and run-v2 with metrics



NAME	BEST_VERSION	EVAL/ACCUR	EVAL/F1	EVAL/LOSS	EVAL/RUNTIME	EVAL/SAMPLE	EVAL/STEPS	HUGGINGFACE
push-model-v1	v1	-	-	-	-	-	-	https://huggingf
run-ver2	-	0.94118	0.94121	0.3575	14.4821	524.786	4.143	-
run-ver1	-	0.94342	0.94342	0.36794	15.0542	504.842	7.905	-


Figure 9: W&B comparison chart — eval/accuracy and eval/loss across both runs

2.2 Task 5 — Push Model to HuggingFace Hub [5 marks]

Version 1 pushed to HuggingFace Hub via `trainer.push_to_hub()`. Model is public at:

<https://huggingface.co/YuvarajK-g25ait2054/ag-news-distilbert>

Model card (MODEL_CARD.md) covers description, hyperparameters, evaluation results, and usage.



File	Size	Version	Commit	Author	Time
main	ag-news-distilbert	269 MB	76a7912	VERIFIED	1 day ago
.gitattributes	1.52 kB	initial commit			3 days ago
README.md	7.37 kB	Update README.md			1 day ago
config.json	853 Bytes	Upload DistilBertForSequenceClassificati...			3 days ago
model.safetensors	268 MB	Upload DistilBertForSequenceClassificati...			3 days ago
tokenizer.json	712 kB	Upload tokenizer			3 days ago
tokenizer_config.json	322 Bytes	Upload tokenizer			3 days ago

Figure 10: HuggingFace model page — YuvarajK-g25ait2054/ag-news-distilbert

3. Docker Deployment and GitHub Actions CI/CD (Tasks 6–7)

3.1 Task 6 — Dockerfile [10 marks]

Docker image hosted at [yuvarajkg25ait2054/ag-news-classifier:latest](https://hub.docker.com/r/yuvarajkg25ait2054/ag-news-classifier:latest).

- Base: `python:3.10-slim`; PyTorch 2.5.1+cpu from PyTorch's CPU wheel index.
- Model pre-baked into image during `docker build` — no internet required at runtime.
- Default CMD runs `--demo` mode: classifies all 4 sample texts from `inference.ipynb` Cell 7.

Demo mode (default) -- runs all 4 notebook texts:

```
docker run --rm yuvarajkg25ait2054/ag-news-classifier:latest
```

Single text:

```
docker run --rm yuvarajkg25ait2054/ag-news-classifier:latest \
```

```
--text "Stock market rises on strong earnings"
```




Figure 11: Docker Hub — `yuvarajkg25ait2054/ag-news-classifier` (publicly pullable)




Figure 12: Docker container — demo mode classifying all 4 sample texts




Figure 13: Docker container — terminal output showing inference results for all 4 categories

3.2 Task 7 — GitHub Actions Workflows [15 marks]

Two workflows under `.github/workflows/`:

- **ci.yml** — triggered on push/PR to `develop` and `main`. Validates required files, runs Python syntax check (`py_compile`), and validates JSON files.
- **inference.yml** — installs dependencies and runs `inference.py` end-to-end with a hard-coded test sentence against the HuggingFace model.

Both workflows pass with green checkmarks.




Figure 14: GitHub Actions — CI Pipeline and Inference workflow with green checkmarks




Figure 15: Inference workflow run log — classification output

4. W&B Experiment Comparison and Conclusion (Task 8)

4.1 Task 8 — W&B Experiment Comparison [5 marks]

The W&B project is publicly accessible at:

<https://wandb.ai/g25ait2054-iit-jodhpur/mlops-assignment3>

The project contains runs **run-v1**, **run-v2**, and **push-model-v1** (logs HuggingFace model URL to W&B summary).

- Version 1 ($lr=2e-5$, $bs=16$) converged smoothly; Version 2 ($lr=5e-5$, $bs=64$) had lower eval loss but slightly worse F1.
- Accuracy difference is small (0.22%); Version 1 selected for deployment based on higher F1.




Figure 16: W&B runs table — run-v1 and run-v2 with accuracy, F1, and loss




Figure 17: W&B comparison view — eval/accuracy and eval/loss with both runs overlaid

4.2 Challenges

- **Docker PyTorch:** Default PyPI `torch==2.5.1` has no CPU wheel; fixed by using PyTorch's CPU index URL (`2.5.1+cpu`).
- **DistilBERT `token_type_ids`:** DistilBERT does not use segment embeddings; had to remove `token_type_ids` from tokenizer output before model call.
- **GitHub Actions inputs:** `github.event.inputs` is empty on push triggers; fixed by hardcoding test values in the workflow YAML.
- **Merge conflicts:** Resolved by keeping `develop` versions of conflicted files.

4.3 Summary of Deliverables

Task	Description	Status	Link / Artefact
Task 1	GitHub Repo + Branching	Done	GitHub
Task 2	Data Prep + Normalisation	Done	<code>prepare_data.ipynb</code>
Task 3	Model Selection	Done	<code>distilbert-base-uncased</code>
Task 4	Training (v1 + v2) + W&B	Done	Kaggle
Task 5	HuggingFace Hub Push	Done	HF Hub
Task 6	Docker Image	Done	Docker Hub
Task 7	GitHub Actions (CI + Infer)	Done	Actions
Task 8	W&B Comparison	Done	W&B